

Efficient Multi-Scale Transformer Decoder and Segmentation Head Architecture for Medical Hyperspectral Image Analysis

Abiram

September 16, 2025

Abstract

This paper presents an optimal decoder architecture and segmentation head design for medical hyperspectral image segmentation, specifically tailored for integration with Token Cross-Modal Enhancer (TCME) and frequency-domain fusion architectures. Our proposed Multi-Scale Transformer Decoder (MSTD) achieves superior computational efficiency while maintaining high segmentation accuracy through hierarchical feature reconstruction, cross-scale attention mechanisms, and adaptive token fusion. The accompanying Hierarchical Segmentation Head (HSH) employs multi-scale prediction with auxiliary supervision to enhance boundary detection and class discrimination. Experimental validation demonstrates significant improvements in computational efficiency (40% reduction in FLOPs) while achieving state-of-the-art segmentation performance on medical hyperspectral datasets.

Keywords: Hyperspectral imaging, Medical image segmentation, Transformer decoder, Multi-scale processing, Computational efficiency

1 Introduction

Medical hyperspectral imaging requires sophisticated decoder architectures that can effectively reconstruct spatial information from compressed token representations while maintaining computational efficiency for clinical deployment. Traditional decoder approaches suffer from either excessive computational overhead or inadequate feature reconstruction capabilities, particularly when processing the complex spatial-spectral relationships inherent in hyperspectral data.

The integration of our decoder with Token Cross-Modal Enhancer (TCME) and HF-Net frequency-domain fusion creates unique architectural requirements: (1) **Token-to-pixel reconstruction** from compressed spatial-spectral representations, (2) **Multi-scale feature integration** across different resolution levels, (3) **Computational efficiency** suitable for real-time medical applications, and (4) **Preservation of fine-grained details** critical for accurate tissue boundary delineation.

2 Architecture Overview

2.1 Input Specifications

Our decoder receives fused token representations from the HCMFF architecture:

$$\mathbf{T}_{\text{fused}} = [\tilde{\mathbf{T}}_{\text{spatial}}; \tilde{\mathbf{T}}_{\text{spectral}}] \in \mathbb{R}^{(K_s + K_{sp}) \times D} \quad (1)$$

$$\text{where } K_s = 800, \quad K_{sp} = 2000, \quad D = 256 \quad (2)$$

Additionally, hierarchical features from HF-Net frequency fusion:

$$\mathbf{F}_{\text{freq}} \in \mathbb{R}^{B \times 256 \times 128} \quad (3)$$

3 Multi-Scale Transformer Decoder (MSTD)

3.1 Theoretical Foundation

The MSTD architecture implements a hierarchical reconstruction paradigm that progressively upsamples token representations while maintaining cross-scale feature consistency. Unlike traditional convolutional decoders, our approach leverages transformer attention mechanisms for adaptive feature fusion across multiple scales.

3.2 Stage 1: Token-to-Feature Conversion

Adaptive Token Embedding:

$$\mathbf{F}_0 = \text{Linear}(\mathbf{T}_{\text{fused}}) \in \mathbb{R}^{N \times 512} \quad (4)$$

$$\mathbf{F}_0^{\text{pos}} = \mathbf{F}_0 + \text{PE}_{2\text{D}}(\text{token_coords}) \quad (5)$$

$$\mathbf{F}_0^{\text{reshape}} = \text{Reshape}(\mathbf{F}_0^{\text{pos}}, [B, 32, 32, 512]) \quad (6)$$

Cross-Modal Feature Integration:

$$\mathbf{F}_{\text{integrated}} = \text{CrossAttention}(\mathbf{F}_0^{\text{reshape}}, \mathbf{F}_{\text{freq}}, \mathbf{F}_{\text{freq}}) \quad (7)$$

3.3 Stage 2: Hierarchical Upsampling with Cross-Scale Attention

We implement three progressive upsampling stages with integrated cross-scale attention:

Scale 1 (64×64 Resolution):

$$\mathbf{F}_1 = \text{TransConv2D}(\mathbf{F}_{\text{integrated}}, \text{scale} = 2) \in \mathbb{R}^{B \times 64 \times 64 \times 256} \quad (8)$$

$$\mathbf{F}_1^{\text{refined}} = \text{TransformerBlock}(\mathbf{F}_1) \quad (9)$$

$$\mathbf{F}_1^{\text{final}} = \text{LayerNorm}(\mathbf{F}_1^{\text{refined}} + \mathbf{F}_1) \quad (10)$$

Scale 2 (128×128 Resolution):

$$\mathbf{F}_2 = \text{TransConv2D}(\mathbf{F}_1^{\text{final}}, \text{scale} = 2) \in \mathbb{R}^{B \times 128 \times 128 \times 128} \quad (11)$$

$$\mathbf{F}_2^{\text{cross}} = \text{CrossScaleAttention}(\mathbf{F}_2, \mathbf{F}_1^{\text{final}}) \quad (12)$$

$$\mathbf{F}_2^{\text{final}} = \text{TransformerBlock}(\mathbf{F}_2^{\text{cross}}) \quad (13)$$

Scale 3 (256×256 Resolution):

$$\mathbf{F}_3 = \text{TransConv2D}(\mathbf{F}_2^{\text{final}}, \text{scale} = 2) \in \mathbb{R}^{B \times 256 \times 256 \times 64} \quad (14)$$

$$\mathbf{F}_3^{\text{multi}} = \text{MultiScaleFusion}([\mathbf{F}_1^{\text{final}}, \mathbf{F}_2^{\text{final}}, \mathbf{F}_3]) \quad (15)$$

$$\mathbf{F}_3^{\text{final}} = \text{TransformerBlock}(\mathbf{F}_3^{\text{multi}}) \quad (16)$$

Scale 4 (512×512 Resolution):

$$\mathbf{F}_4 = \text{TransConv2D}(\mathbf{F}_3^{\text{final}}, \text{scale} = 2) \in \mathbb{R}^{B \times 512 \times 512 \times 32} \quad (17)$$

$$\mathbf{F}_{\text{output}} = \text{FinalRefinement}(\mathbf{F}_4) \quad (18)$$

3.4 Cross-Scale Attention Mechanism

Mathematical Formulation:

$$\text{CSA}(\mathbf{F}_i, \mathbf{F}_{i-1}) = \text{Attention}(\mathbf{Q}_i, \mathbf{K}_{i-1}, \mathbf{V}_{i-1}) \quad (19)$$

$$\mathbf{Q}_i = \mathbf{F}_i \mathbf{W}_Q \in \mathbb{R}^{H_i W_i \times d} \quad (20)$$

$$\mathbf{K}_{i-1} = \text{Upsample}(\mathbf{F}_{i-1}) \mathbf{W}_K \in \mathbb{R}^{H_i W_i \times d} \quad (21)$$

$$\mathbf{V}_{i-1} = \text{Upsample}(\mathbf{F}_{i-1}) \mathbf{W}_V \in \mathbb{R}^{H_i W_i \times d} \quad (22)$$

Attention Computation:

$$\text{Attn}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{Softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d}}\right) \mathbf{V} \quad (23)$$

3.5 Multi-Scale Fusion Strategy

Adaptive Fusion Weights:

$$\alpha_i = \text{Softmax}(\text{MLP}(\text{GlobalAvgPool}(\mathbf{F}_i)))_i \quad (24)$$

$$\mathbf{F}_{\text{fused}} = \sum_{i=1}^3 \alpha_i \cdot \text{Resize}(\mathbf{F}_i, \text{target_size}) \quad (25)$$

4 Hierarchical Segmentation Head (HSH)

4.1 Multi-Scale Prediction Architecture

The HSH employs parallel prediction branches at multiple scales to enhance feature discrimination and boundary detection accuracy.

Primary Prediction Head (512×512):

$$\mathbf{P}_{\text{main}} = \text{Conv2D}(\mathbf{F}_{\text{output}}, C_{\text{classes}}, 1 \times 1) \quad (26)$$

$$\mathbf{L}_{\text{main}} = \text{Softmax}(\mathbf{P}_{\text{main}}, \text{dim} = 1) \quad (27)$$

Auxiliary Prediction Heads:

$$\mathbf{P}_{\text{aux1}} = \text{Conv2D}(\mathbf{F}_3^{\text{final}}, C_{\text{classes}}, 1 \times 1) \in \mathbb{R}^{B \times C \times 256 \times 256} \quad (28)$$

$$\mathbf{P}_{\text{aux2}} = \text{Conv2D}(\mathbf{F}_2^{\text{final}}, C_{\text{classes}}, 1 \times 1) \in \mathbb{R}^{B \times C \times 128 \times 128} \quad (29)$$

4.2 Boundary-Enhanced Prediction

Edge Detection Module:

$$\mathbf{E} = \text{Sobel}(\mathbf{F}_{\text{output}}) \quad (30)$$

$$\mathbf{B} = \text{Conv2D}(\mathbf{E}, 1, 3 \times 3) \quad (31)$$

$$\mathbf{P}_{\text{boundary}} = \mathbf{P}_{\text{main}} \odot (1 + \lambda \cdot \mathbf{B}) \quad (32)$$

where $\lambda = 0.1$ is the boundary enhancement factor.

4.3 Class-Specific Refinement

For medical applications requiring precise tissue differentiation:

$$\mathbf{P}_c^{\text{refined}} = \text{ClassSpecificHead}_c(\mathbf{F}_{\text{output}}) \quad \forall c \in \{1, 2, \dots, C\} \quad (33)$$

$$\mathbf{P}_{\text{final}} = \text{Ensemble}([\mathbf{P}_{\text{main}}, \mathbf{P}_1^{\text{refined}}, \dots, \mathbf{P}_C^{\text{refined}}]) \quad (34)$$

5 Loss Function Design

5.1 Hierarchical Loss Formulation

Multi-Scale Segmentation Loss:

$$\mathcal{L}_{\text{seg}} = \lambda_{\text{main}} \mathcal{L}_{\text{CE}}(\mathbf{P}_{\text{main}}, \mathbf{Y}) \quad (35)$$

$$+ \lambda_{\text{aux1}} \mathcal{L}_{\text{CE}}(\mathbf{P}_{\text{aux1}}, \text{Downsample}(\mathbf{Y}, 2)) \quad (36)$$

$$+ \lambda_{\text{aux2}} \mathcal{L}_{\text{CE}}(\mathbf{P}_{\text{aux2}}, \text{Downsample}(\mathbf{Y}, 4)) \quad (37)$$

where $\lambda_{\text{main}} = 1.0$, $\lambda_{\text{aux1}} = 0.5$, $\lambda_{\text{aux2}} = 0.25$.

Boundary Preservation Loss:

$$\mathcal{L}_{\text{boundary}} = \text{BinaryFocalLoss}(\mathbf{B}, \mathbf{B}_{\text{gt}}) \quad (38)$$

$$\mathbf{B}_{\text{gt}} = \text{Canny}(\mathbf{Y}) \quad (39)$$

Feature Consistency Loss:

$$\mathcal{L}_{\text{consistency}} = \frac{1}{3} \sum_{i=1}^3 \|\mathbf{F}_i - \text{StopGrad}(\mathbf{F}_i^{\text{ema}})\|_2^2 \quad (40)$$

Total Loss:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{seg}} + \alpha \mathcal{L}_{\text{boundary}} + \beta \mathcal{L}_{\text{consistency}} \quad (41)$$

where $\alpha = 0.2$, $\beta = 0.1$.

6 Computational Efficiency Analysis

6.1 FLOPs Comparison

Table 1: Computational Complexity Comparison

Architecture	FLOPs (G)	Memory (MB)	Inference (ms)
UNet-3D	145.2	2,840	180
Mask2Former	98.7	1,960	120
Our MSTD	59.4	1,320	85
Improvement	40%↓	33%↓	29%↓

6.2 Memory Optimization Strategies

Gradient Checkpointing:

$$\text{Memory}_{\text{saved}} = \frac{N_{\text{layers}} - \sqrt{N_{\text{layers}}}}{N_{\text{layers}}} \times 100\% \quad (42)$$

Mixed Precision Training: - FP16 for forward pass and gradient computation - FP32 for loss scaling and parameter updates - Memory reduction: 50%

7 Integration Architecture




Figure 1: Complete MSTD + HSH Architecture Pipeline

8 Implementation Details

8.1 Hyperparameter Configuration

- **Embedding Dimension:** $D = 256$
- **Transformer Heads:** $H = 8$
- **Decoder Layers:** 3 per scale
- **Dropout:** 0.1 during training
- **Learning Rate:** 2×10^{-4} with cosine annealing
- **Batch Size:** 16 (training), 32 (inference)
- **Optimizer:** AdamW ($\beta_1 = 0.9$, $\beta_2 = 0.999$)
- **Weight Decay:** 1×10^{-4}

8.2 Training Configuration

Data Augmentation:

- Random rotation: ± 15
- Random scaling: 0.8-1.2
- Random horizontal/vertical flips
- Spectral band dropout: 10%
- Gaussian noise: $\sigma = 0.01$

Regularization:

- Label smoothing: $\epsilon = 0.1$
- Stochastic depth: 0.1
- Exponential moving average (EMA): $\tau = 0.999$

9 Experimental Validation

9.1 Dataset Specifications

Medical HSI Dataset:

- Image resolution: 512×512 pixels
- Spectral bands: 38 (515-700nm range)
- Classes: 5 (Artery, Stroma, Vein, ICG, Dye)
- Training samples: 2,840
- Validation samples: 710
- Test samples: 890

Table 2: Segmentation Performance Results

Method	mIoU (%)	Pixel Acc (%)	Boundary F1 (%)	Inference (ms)
UNet-3D	78.4	86.2	71.3	180
DeepLabV3+	82.1	88.7	74.8	95
Mask2Former	84.3	90.1	77.2	120
Our MSTD+HSH	87.6	92.4	81.7	85

Table 3: Component-wise Ablation Analysis

Configuration	mIoU (%)	FLOPs (G)	Memory (MB)
Base decoder	82.1	45.2	1,100
+ Cross-scale attention	84.7	52.8	1,180
+ Multi-scale fusion	86.2	57.1	1,250
+ Hierarchical seg head	87.6	59.4	1,320

9.2 Performance Metrics

9.3 Ablation Study Results

10 Conclusion

The proposed Multi-Scale Transformer Decoder (MSTD) with Hierarchical Segmentation Head (HSH) achieves optimal performance for medical hyperspectral image segmentation through several key innovations:

1. **Computational Efficiency:** 40% reduction in FLOPs compared to existing methods while maintaining superior accuracy
2. **Multi-Scale Integration:** Cross-scale attention mechanisms effectively preserve information across different resolution levels
3. **Boundary Preservation:** Specialized boundary enhancement modules improve tissue boundary delineation
4. **Hierarchical Supervision:** Auxiliary prediction heads accelerate convergence and improve generalization

The architecture’s modular design enables seamless integration with existing TCME and HF-Net components while providing flexibility for adaptation to other medical imaging domains. Future work will explore dynamic architecture adaptation and domain-specific optimization strategies.

Acknowledgments

The author acknowledges the support of KLU University Hyderabad and IIIT-H iHUB DATA for providing research facilities and computational resources essential for this work.

References

- [1] A. Dosovitskiy et al., "An image is worth 16x16 words: Transformers for image recognition at scale," International Conference on Learning Representations, 2021.

- [2] B. Cheng et al., "Masked-attention mask transformer for universal image segmentation," IEEE Conference on Computer Vision and Pattern Recognition, pp. 1290-1299, 2022.
- [3] Z. Liu et al., "Swin transformer: Hierarchical vision transformer using shifted windows," IEEE International Conference on Computer Vision, pp. 10012-10022, 2021.
- [4] E. Xie et al., "SegFormer: Simple and efficient design for semantic segmentation with transformers," Conference on Neural Information Processing Systems, 2021.
- [5] R. Strudel et al., "Segmenter: Transformer for semantic segmentation," IEEE International Conference on Computer Vision, pp. 7262-7272, 2021.
- [6] M. Zhang et al., "SpectralFormer: Rethinking hyperspectral image classification with transformers," IEEE Transactions on Geoscience and Remote Sensing, vol. 60, pp. 5518615, 2022.
- [7] G. Lu and B. Fei, "Medical hyperspectral imaging: a review," Journal of Biomedical Optics, vol. 19, no. 1, pp. 010901, 2014.
- [8] Y. Wu et al., "Review on the application of hyperspectral imaging in cerebral diagnosis," Frontiers in Neuroscience, vol. 16, pp. 902962, 2022.