Enhanced Hierarchical Frequency-Domain Fusion Architecture for Medical Hyperspectral Imaging: A Novel Multi-Scale Cross-Modal Integration Approach

Abiram

September 2025

Abstract

Medical hyperspectral imaging (HSI) presents unique challenges in effectively fusing spatial and spectral information for accurate tissue characterization and disease detection. This paper proposes an enhanced hierarchical frequency-domain fusion architecture, termed HF-Net, that addresses fundamental limitations in existing cross-modal fusion approaches. Our novel contribution lies in the development of a sophisticated multi-scale frequency processing mechanism that intelligently combines four frequency components—fused amplitude, aligned phase, low-frequency global structure, and high-frequency fine details—through adaptive octave-scale processing. Unlike conventional spatial-domain fusion methods, our approach leverages frequency-domain properties to achieve superior noise robustness, computational efficiency, and preservation of critical spectral-spatial relationships. The architecture incorporates a two-stage fusion strategy: cross-modal frequency preparation followed by hierarchical multi-scale integration with learnable attention weights. Experimental validation demonstrates significant improvements in segmentation accuracy while maintaining real-time processing capabilities essential for clinical applications. The proposed method establishes a new paradigm for medical HSI analysis by providing theoretically grounded and practically validated solutions for complex multi-modal medical imaging challenges.

1 Introduction

Hyperspectral imaging has emerged as a revolutionary technology in medical diagnostics, offering unprecedented capabilities for non-invasive tissue characterization and pathological detection [1]. The technology's ability to capture rich spatial and spectral information simultaneously provides clinicians with detailed insights into tissue composition, oxygenation levels, and metabolic activity that are invisible to conventional imaging modalities [2].

Current medical HSI applications span diverse clinical domains, from intraoperative brain tumor detection and surgical guidance to dermatological lesion analysis and ophthalmologic diagnostics [3]. However, the inherent complexity of medical HSI data presents significant computational and methodological challenges that existing approaches have not adequately addressed.

1.1 Challenges in Medical HSI Processing

Medical hyperspectral data processing faces four fundamental challenges: (1) **Cross-modal information fusion**—effectively integrating spatial morphological features with spectral biochemical signatures; (2) **Noise sensitivity**—managing the high noise-to-signal ratio characteristic of medical imaging environments; (3) **Real-time processing requirements**—achieving clinically acceptable processing speeds for intraoperative applications; and (4) **Preservation of critical diagnostic information**—maintaining both fine anatomical details and broad contextual relationships.

Traditional approaches to medical HSI processing typically operate in the spatial domain, employing convolutional neural networks or attention mechanisms to extract and fuse features [5, 6]. While these methods have shown promise, they suffer from computational inefficiency and limited ability to handle the unique frequency characteristics inherent in hyperspectral data.

1.2 Limitations of Existing Approaches

Current state-of-the-art methods in medical HSI processing can be categorized into three main approaches:

Spatial-domain methods rely on traditional CNN architectures with attention mechanisms to fuse spatial and spectral information. However, these approaches exhibit $O(N^2)$ computational complexity for global context modeling and struggle with preserving spectral relationships during feature extraction [4].

Spectral-domain methods focus on utilizing spectral signatures for classification and segmentation but often neglect spatial contextual information crucial for accurate medical diagnosis [7].

Hybrid approaches attempt to combine spatial and spectral processing but lack principled frameworks for optimal information fusion and typically require extensive computational resources [8].

1.3 Our Novel Contributions

This work introduces HF-Net (Hierarchical Frequency Network), a novel architecture that addresses these limitations through principled frequency-domain processing. Our key contributions include:

1. **Frequency-domain fusion paradigm**: A novel approach that transforms both spatial and spectral information into the frequency domain for optimal cross-modal integration.

2. **Hierarchical multi-scale processing**: An innovative octave-scale processing mechanism that adaptively balances global anatomical context with fine pathological details through learnable frequency component weighting.

3. **Four-component fusion strategy**: A sophisticated integration of fused amplitude, aligned phase, low-frequency structure, and high-frequency details through dynamic scaling and attention mechanisms.

4. **Computational efficiency**: Achievement of $O(N \log N)$ complexity through FFT-based processing while maintaining superior segmentation accuracy.

The remainder of this paper is organized as follows: Section 2 presents a comprehensive literature review and positions our work within the current research landscape. Section 3 details the mathematical foundations and architectural design of HF-Net. Section 4 provides experimental validation and comparative analysis. Section 5 discusses implications and future directions, and Section 6 concludes the work.

2 Literature Review and Related Work

2.1 Medical Hyperspectral Imaging Fundamentals

Medical hyperspectral imaging technology has evolved significantly since its initial applications in remote sensing [1]. The fundamental principle involves capturing hundreds of contiguous spectral bands across the electromagnetic spectrum, typically in the visible and near-infrared (VNIR) range of 400-1000 nm, where biological chromophores exhibit distinct absorption characteristics.

Recent comprehensive reviews [4,10] have highlighted the technology's potential across various medical specialties. In neurosurgery, HSI enables real-time identification of tumor boundaries by detecting changes in tissue oxygenation and metabolic activity [2]. Dermatological applications leverage spectral signatures to differentiate between benign and malignant skin lesions. Ophthalmologic uses include retinal vessel analysis and early diabetic retinopathy detection.

2.2 Frequency-Domain Processing in Medical Imaging

The application of frequency-domain techniques in medical imaging has demonstrated significant advantages, particularly in denoising and feature enhancement applications [8]. Fast Fourier Transform (FFT) based methods offer natural noise separation capabilities and efficient computation of global image statistics.

Recent work by Duan et al. [7] demonstrated the potential of frequency-domain processing for spectral super-resolution in hyperspectral imaging. Their approach showed that frequency-domain feature learning could capture spectral-spatial correlations more effectively than traditional spatial-domain methods. However, their work focused primarily on reconstruction tasks rather than the complex fusion requirements of medical applications.

2.3 Multi-Scale Hierarchical Processing

Multi-scale processing has become a cornerstone of modern medical image analysis, with several notable architectures emerging in recent years [5,9]. The Dynamic Hierarchical Multi-Scale Fusion Network (DHMF-MLP) proposed by Cheng et al. [9] demonstrated the effectiveness of hierarchical feature fusion for medical image segmentation, achieving significant improvements in small lesion detection.

HiFuse [5] introduced a three-branch hierarchical architecture that effectively combines CNN and Transformer features across multiple scales. Their approach achieved notable performance improvements on various medical imaging datasets, highlighting the importance of multi-scale feature integration in medical applications.

HiFormer [6] further advanced the field by proposing hierarchical multi-scale representations using transformers specifically for medical image segmentation, demonstrating superior performance through efficient CNN-Transformer bridging.

2.4 Cross-Modal Fusion in Medical Imaging

Cross-modal fusion represents a critical challenge in medical imaging, where different modalities or feature types must be effectively integrated. Recent work has explored various fusion strategies, from simple concatenation to sophisticated attention mechanisms [11].

The development of spectral-spatial fusion methods has been particularly active in the hyperspectral imaging community. However, most existing approaches operate exclusively in the spatial domain, missing opportunities for more efficient and effective frequency-domain integration.

2.5 Gaps in Current Research

Despite significant advances, several critical gaps remain in medical HSI processing:

1. **Limited frequency-domain exploration**: Few works have systematically explored frequency-domain processing for medical HSI applications.

2. **Insufficient cross-modal integration**: Existing methods lack principled approaches to optimally combine spatial morphological and spectral biochemical information.

3. **Scale-dependent processing limitations**: Current multi-scale methods do not adequately address the unique characteristics of different frequency scales in medical data.

4. **Computational efficiency constraints**: Many state-of-the-art methods require excessive computational resources, limiting their clinical applicability.

3 Methodology: HF-Net Architecture

3.1 Theoretical Foundation

The HF-Net architecture is grounded in fundamental frequency-domain signal processing principles adapted for medical hyperspectral imaging. The core insight is that spatial and spectral information can be more effectively fused in the frequency domain, where noise separation, global context modeling, and cross-modal correlation analysis can be performed more efficiently.

For hyperspectral data $H(x, y, \lambda)$ with spatial dimensions (x, y) and spectral dimension λ, we define two complementary frequency domain transformations:

Spatial Frequency Transform:

$$F_s(u, v, \lambda) = \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} H(x, y, \lambda) \cdot e^{-j2\pi(ux/M+vy/N)} \tag{1}$$

Spectral Frequency Transform:

$$F_\lambda(x, y, f) = \sum_{\lambda=0}^{L-1} H(x, y, \lambda) \cdot e^{-j2\pi f\lambda/L} \tag{2}$$

where (u, v) represent spatial frequency coordinates, f represents spectral frequency, and M, N, L denote the spatial and spectral dimensions respectively.

3.2 Architecture Overview

The HF-Net architecture consists of five sequential stages that systematically transform, fuse, and reconstruct spatial-spectral information:

1. **Dual-Domain Feature Extraction**: Independent processing of spatial and spectral tokens

2. **Frequency Domain Transformation**: Conversion to frequency representation with amplitude-phase separation

3. **Cross-Modal Frequency Fusion**: Four-component preparation stage

4. **Hierarchical Multi-Scale Processing**: Octave-scale fusion with adaptive weighting

5. **Inverse Transform & Reconstruction**: Spatial domain reconstruction for downstream processing

3.3 Stage 1: Dual-Domain Feature Extraction

The architecture begins with parallel processing of spatial and spectral information using specialized encoders designed to preserve domain-specific characteristics.

Spatial Branch Processing:

$$\mathbf{F}_s^{enh} = \text{SpatialEncoder}(\mathbf{T}_s) \tag{3}$$

$$: \mathbb{R}^{B\times256\times64} \rightarrow \mathbb{R}^{B\times256\times128} \tag{4}$$

where $\mathbf{T}_s$ represents spatial tokens extracted from MT-SCFormer processing, and B denotes the batch size.

Spectral Branch Processing:

$$\mathbf{F}_\lambda^{enh} = \text{SpectralEncoder}(\mathbf{T}_\lambda) \tag{5}$$

$$: \mathbb{R}^{B\times224\times64} \to \mathbb{R}^{B\times224\times128} \tag{6}$$

where $\mathbf{T}_\lambda$ represents spectral tokens from 3D-Mamba processing.

3.4 Stage 2: Frequency Domain Transformation

Each enhanced feature set undergoes domain-specific frequency transformation to extract amplitude and phase components essential for subsequent fusion.

Spatial-to-Frequency Conversion:

$$\mathbf{F}_s^{2D} = \text{Reshape}(\mathbf{F}_s^{enh}, 16, 16, 128) \tag{7}$$

$$\mathbf{F}_s^{freq} = \text{FFT2D}(\mathbf{F}_s^{2D}, \text{dims} = (1, 2)) \tag{8}$$

$$A_s = |\mathbf{F}_s^{freq}| \tag{9}$$

$$P_s = \angle\mathbf{F}_s^{freq} \tag{10}$$

Spectral-to-Frequency Conversion:

$$\mathbf{F}_\lambda^{freq} = \text{FFT1D}(\mathbf{F}_\lambda^{enh}, \text{dim} = -1) \tag{11}$$

$$A_\lambda = |\mathbf{F}_\lambda^{freq}| \tag{12}$$

$$P_\lambda = \angle\mathbf{F}_\lambda^{freq} \tag{13}$$

3.5 Stage 3: Cross-Modal Frequency Fusion

This stage represents a critical innovation in our approach, where four distinct frequency components are prepared for hierarchical processing. Unlike traditional fusion methods that operate directly on features, our approach creates specialized frequency representations optimized for different aspects of medical image analysis.

Component 1: Fused Amplitude

$$A_f = w_s A_s + w_\lambda A_\lambda \tag{14}$$

where w_s and w_λ are learnable parameters that balance spatial morphological and spectral biochemical contributions.

Component 2: Aligned Phase Phase alignment is achieved through multi-head attention that learns optimal phase relationships between spatial structural and spectral compositional information:

$$\mathbf{P}_{concat} = \text{Concat}([P_s, P_\lambda], \text{dim} = -1) \tag{15}$$

$$P_f = \text{MultiHeadAttention}(\mathbf{P}_{concat}, \mathbf{P}_{concat}, \mathbf{P}_{concat}) \tag{16}$$

Component 3: Low-Frequency Structure

$$F_{low} = \text{Conv1D}_{LowPass}(A_f) \tag{17}$$

This component captures global anatomical structures and organ boundaries essential for medical diagnosis.

Component 4: High-Frequency Details

$$F_{high} = \text{Conv1D}_{HighPass}(A_f) \tag{18}$$

This component preserves fine pathological details and tissue texture variations crucial for accurate lesion detection.

3.6 Stage 4: Hierarchical Multi-Scale Processing - The Core Innovation

The hierarchical multi-scale processing stage represents our primary methodological contribution. This stage implements adaptive octave-scale processing that intelligently combines the four frequency components at different scales to balance global anatomical context with fine pathological details.

3.6.1 Octave-Scale Component Integration

For each octave scale $s \in \{1, 2, 4, 8\}$, we perform scale-specific component fusion:

Scale 1 (Full Resolution):

$$A_1^{combined} = F_{low} + F_{high} \tag{19}$$

At full resolution, both global structure and fine details are equally preserved.

Scale 2 (Half Resolution):

$$F_{low}^{(2)} = \text{AvgPool1D}(F_{low}, \text{kernel} = 2) \tag{20}$$
$$F_{high}^{(2)} = \text{AvgPool1D}(F_{high}, \text{kernel} = 2) \times 0.5 \tag{21}$$
$$A_2^{combined} = F_{low}^{(2)} + F_{high}^{(2)} \tag{22}$$

High-frequency components are attenuated to emphasize global patterns while maintaining some detail preservation.

Scale 4 (Quarter Resolution):

$$F_{low}^{(4)} = \text{AvgPool1D}(F_{low}, \text{kernel} = 4) \tag{23}$$
$$F_{high}^{(4)} = \text{AvgPool1D}(F_{high}, \text{kernel} = 4) \times 0.25 \tag{24}$$
$$A_4^{combined} = F_{low}^{(4)} + F_{high}^{(4)} \tag{25}$$

Further high-frequency attenuation focuses on broader anatomical structures.

Scale 8 (Eighth Resolution):

$$A_8^{combined} = \text{AvgPool1D}(F_{low}, \text{kernel} = 8) \tag{26}$$

Only low-frequency global structure is retained for coarsest-scale processing.

3.6.2 Scale-Specific Processing

Each combined amplitude undergoes dedicated processing optimized for its resolution:

$$\mathbf{F}_s^{processed} = \text{ScaleProcessor}_s(A_s^{combined}) \tag{27}$$
$$\mathbf{F}_s^{upsampled} = \text{Interpolate}(\mathbf{F}_s^{processed}, \text{size} = A_f.\text{shape}[-1]) \tag{28}$$

3.6.3 Adaptive Multi-Scale Fusion

The processed features from all scales are combined using learned attention weights:

$$\boldsymbol{\alpha} = \text{Softmax}(\mathbf{W}_{scale}) \tag{29}$$
$$A_{final} = \sum_s \alpha_s \mathbf{F}_s^{upsampled} \tag{30}$$

where $\mathbf{W}_{scale}$ represents learnable scale importance weights that adapt to different medical imaging scenarios.

3.7 Stage 5: Inverse Transform & Reconstruction

The final stage reconstructs spatial-domain features by combining the processed amplitude with aligned phase information:

$$\mathbf{F}_{complex} = A_{final} \cdot e^{jP_f} \tag{31}$$

$$\mathbf{F}_{spatial} = \text{Real}(\text{IFFT2D}(\mathbf{F}_{complex})) \tag{32}$$

The resulting spatial features $\mathbf{F}_{spatial}$ preserve both global anatomical context and fine pathological details while benefiting from frequency-domain noise suppression and efficient processing.

3.8 Theoretical Advantages

The HF-Net architecture provides several theoretical advantages over existing approaches:

1. **Computational Efficiency**: O(N log N) complexity through FFT-based processing versus O(N^2) for attention-based spatial methods.
2. **Noise Robustness**: Natural noise separation in frequency domain provides superior performance in noisy medical imaging environments.
3. **Global Context**: Instantaneous global receptive field through FFT transformation eliminates the need for deep architectures to capture long-range dependencies.
4. **Spectral Preservation**: Dedicated spectral frequency processing maintains critical biochemical information often lost in spatial-domain approaches.
5. **Adaptive Multi-Scale Processing**: Learnable scale fusion enables automatic adaptation to different medical imaging scenarios and pathological conditions.

4 Experimental Validation and Results

4.1 Dataset and Experimental Setup

The proposed HF-Net architecture was evaluated on multiple medical hyperspectral imaging datasets to demonstrate its effectiveness across different medical applications. The experimental validation focused on segmentation accuracy, computational efficiency, and robustness to noise—critical factors for clinical deployment.

Datasets: Experiments were conducted on publicly available medical HSI datasets including brain tumor hyperspectral images, dermatological lesion data, and synthetic medical HSI datasets with controlled noise characteristics. All datasets were preprocessed using standard normalization and augmentation techniques to ensure fair comparison with baseline methods.

Implementation Details: The architecture was implemented using PyTorch with CUDA acceleration. Training employed the Adam optimizer with initial learning rate 1e-3, batch size 8, and 100 epochs. The scale processors consisted of 3-layer CNNs with ReLU activation and batch normalization.

4.2 Comparative Analysis

Comparisons were performed against state-of-the-art medical image segmentation methods, including traditional CNN approaches, Transformer-based architectures, and existing HSI-specific methods. Evaluation metrics included Dice coefficient, IoU, pixel accuracy, and computational metrics (FLOPs, inference time).

Segmentation Performance: HF-Net achieved superior segmentation accuracy across all evaluated datasets, with particularly notable improvements in small lesion detection and boundary preservation. The frequency-domain processing demonstrated superior noise robustness compared to spatial-domain alternatives.

Computational Efficiency: The O(N log N) complexity of FFT-based processing resulted in significant speed improvements, making the method suitable for real-time clinical applications. The hierarchical multi-scale processing added minimal computational overhead while providing substantial accuracy gains.

5 Discussion and Clinical Implications

5.1 Technical Contributions

The HF-Net architecture introduces several technical innovations that advance the state-of-the-art in medical hyperspectral imaging:

1. **Principled frequency-domain fusion**: The systematic transformation and fusion of spatial-spectral information in the frequency domain provides a theoretically grounded approach to cross-modal integration.

2. **Adaptive multi-scale processing**: The octave-scale processing mechanism automatically balances global anatomical context with fine pathological details through learned attention weights.

3. **Four-component integration strategy**: The systematic preparation and fusion of amplitude, phase, low-frequency, and high-frequency components provides comprehensive representation of medical image characteristics.

5.2 Clinical Relevance

The proposed architecture addresses several critical requirements for clinical deployment:

Real-time processing capability through efficient frequency-domain computation enables intraoperative applications where immediate feedback is essential.

Noise robustness is crucial in clinical environments where imaging conditions may be suboptimal due to patient movement, lighting variations, or equipment limitations.

Preservation of diagnostic information at multiple scales ensures that both global anatomical structures and fine pathological details are maintained throughout processing.

5.3 Limitations and Future Work

While the HF-Net architecture demonstrates significant advantages, several limitations warrant discussion:

Memory requirements for frequency domain transformations may be substantial for very high-resolution hyperspectral data, though this is offset by computational efficiency gains.

Hyperparameter sensitivity in scale weighting and frequency filtering may require domain-specific tuning for optimal performance across different medical applications.

Future work will explore extensions to other medical imaging modalities, investigation of learned frequency filtering approaches, and development of domain adaptation techniques for cross-institutional deployment.

6 Conclusion

This paper introduces HF-Net, a novel hierarchical frequency-domain fusion architecture for medical hyperspectral imaging that addresses fundamental limitations in existing cross-modal fusion approaches. The key innovation lies in the systematic processing of four frequency components—fused amplitude, aligned phase, low-frequency structure, and high-frequency details—through adaptive octave-scale integration.

The theoretical foundation demonstrates superior computational efficiency (O(N log N) complexity), enhanced noise robustness through frequency-domain processing, and preservation of critical medical imaging characteristics through multi-scale hierarchical fusion. Experimental validation confirms significant improvements in segmentation accuracy while maintaining computational efficiency essential for clinical applications.

The proposed architecture establishes a new paradigm for medical HSI processing by providing principled, theoretically grounded solutions for complex multi-modal medical imaging challenges. The frequency-domain approach opens new research directions for efficient and effective medical image analysis, with immediate applicability to clinical diagnostic and surgical guidance systems.

Future research directions include extension to multi-modal medical imaging fusion, investigation of learned frequency filtering techniques, and development of domain adaptation methods for cross-institutional deployment. The HF-Net architecture provides a robust foundation for advancing medical hyperspectral imaging from research laboratory to clinical practice.

Acknowledgments

The author acknowledges the support of KLH University Hyderabad and IIIT-H iHUB DATA for providing research facilities and computational resources essential for this work.

References

[1] G. Lu and B. Fei, "Medical hyperspectral imaging: a review," *Journal of Biomedical Optics*, vol. 19, no. 1, pp. 010901, 2014.

[2] Y. Wu, F. Gao, and H. Li, "Review on the application of hyperspectral imaging in cerebral diagnosis," *Frontiers in Neuroscience*, vol. 16, pp. 902962, 2022.

[3] H. Fabelo et al., "An intraoperative visualization system using hyperspectral imaging to aid in brain tumor delineation," *Sensors*, vol. 18, no. 2, pp. 430, 2019.

[4] A. Bhargava and A. Bansal, "Hyperspectral imaging and its applications: A review," *Heliyon*, vol. 10, no. 13, pp. e33208, 2024.

[5] X. Huo et al., "HiFuse: Hierarchical multi-scale feature fusion network for medical image classification," *arXiv preprint arXiv:2209.10218*, 2022.

[6] M. Heidari et al., "HiFormer: Hierarchical multi-scale representations using transformers for medical image segmentation," *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pp. 6202–6212, 2023.

[7] P. Duan et al., "Spectral super-resolution in frequency domain," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 36, no. 7, pp. 12338–12348, 2025.

[8] H. Li and D. Wu, "Hybrid-domain synergistic transformer for hyperspectral image denoising," *arXiv preprint arXiv:2507.20099*, 2025.

[9] Z. Cheng et al., "Dynamic hierarchical multi-scale fusion network with axial mlp for medical image segmentation," *Scientific Reports*, vol. 13, no. 1, pp. 6342, 2023.

[10] A. ul Rehman et al., "A review of the medical hyperspectral imaging systems and applications," *Journal of Biophotonics*, vol. 14, no. 7, pp. e202000360, 2021.

[11] X. Huo et al., "Multi-task hierarchical model of medical images based on alternate interaction," *Computer Methods and Programs in Biomedicine*, vol. 247, pp. 108099, 2025.